

Agmarknet Scraper VM — Command Reference

Quick-reference cheat sheet for operating the scraper on the VM (`sethu-admin1` , Azure, `104.211.91.182`) day to day. For the full deployment story (what was built and why), see `agmarknet_scraper_vm_deployment.md` / `.pdf` in this same folder.

Connecting

```
ssh -i "sethu-admin1_key.pem" azureuser@104.211.91.182
```

Checking scraper status

Is it running right now?

```
sudo systemctl status agmarknet-scraper
```

Look for `Active: active (running)`.

Watch it live (updates in real time; Ctrl+C to stop watching — this does *not* stop the scraper itself):

```
sudo journalctl -u agmarknet-scraper -f
```

Just the progress count (how many commodities done in the current cycle):

```
sudo journalctl -u agmarknet-scraper --no-pager | grep "commodities done" | tail -10
```

Check for errors/failures:

```
sudo journalctl -u agmarknet-scraper --no-pager | grep -iE "error|failed|traceback" | tail
```

Confirm a cycle actually saved to the database (look for the "upserted" line — this only appears once a full cycle finishes):

```
sudo journalctl -u agmarknet-scraper --no-pager | grep "upserted" | tail -5
```

Controlling the service

Restart it manually:

```
sudo systemctl restart agmarknet-scraper
```

Stop it (won't restart until you start it again, even after a reboot):

```
sudo systemctl stop agmarknet-scraper
```

Start it again:

```
sudo systemctl start agmarknet-scraper
```

Disable it from auto-starting on boot (rarely needed — it's meant to stay enabled):

```
sudo systemctl disable agmarknet-scraper
```

Manually backfilling a specific past date

Runs once and exits — not the always-on daemon. Useful if you notice a specific day is missing and don't want to wait for the daemon's own automatic 3-day gap-check:

```
cd ~/projects/pipecat
source venv/bin/activate
python -m bot_processors.pricing.agmarknet_scraper --date=2026-08-09 --resume
```

`--resume` skips any commodity that already has data for that date — safe to re-run if it fails partway.

Checking resource usage

Memory/CPU used by the service:

```
sudo systemctl status agmarknet-scraper --no-pager | grep -E "Memory|CPU"
```

Overall VM memory/disk:

```
free -h
df -h /
```

Updating the code

The VM pulls from GitHub via a read-only deploy key — it can fetch updates but never push:

```
cd ~/projects/pipecat
git pull
sudo systemctl restart agmarknet-scraper
```

Checking the actual scraped data

The VM itself doesn't store the data — every row gets written straight to Postgres on the dev Windows PC over the Tailscale tunnel. To see the real data (not just progress logs), use pgAdmin or a query **on the Windows side**, not the VM:

```
-- one row per day, sorted, with a count
SELECT arrival_date, count(*) FROM fact_market_prices
GROUP BY arrival_date ORDER BY arrival_date;

-- every row, pre-sorted (the fact_market_prices_ordered view)
SELECT * FROM fact_market_prices_ordered;
```

Tailscale (networking)

Check the tunnel is up:

```
sudo tailscale status
```

Check connectivity to the Windows PC specifically:

```
sudo tailscale ping -c 2 100.109.45.19
```